

Sales Forecast, Example

Practice exercise: Seasonal Naive Method, Exponential Smoothing Model and ARIMA Model

JDM

Sales Forecast Considerations:

- This file shows the result of “Sales_per_day_Forecast.R” file. It is the same code.
- Although RMarkdown idea is to hide the script, some important scripts will be shown.

This EXAMPLE shows in a simple way how to use some popular methods for Forecast: Benchmark Method (Seasonal Naive Method), Exponential Smoothing Model and ARIMA Model.


This data is not adjusted for inflation. So these are nominal values. (practice excersice) Monthly retail sales per day

```
# Load data
data <- read.csv("https://raw.githubusercontent.com/Jdmuci/R_practices/main/sales_per_day.csv")

# declare as time series data
Y <- ts(data[,2], start = c(1992,1), end = c(2017,12), frequency = 12)
```

Preliminary analysis

Time plot




1- Seasonality and Trend

Data has a strong trend and seasonality. Need to investigate transformations.

- Take the first difference to remove trend.


```
DY <- diff(Y)
```

Time plot of differenced data:




Series appears trend-stationary, but need to investigate seasonality.

Lets now focus on seasonality. The next two plots are interesting and useful to analyse seasonality.



This function makes a series for each year and plot them together in a 12 months period to see the seasonality. Apparently there is a drop between december and january, and a rise from november to december (christmas season).

Lets look at another seasonal plot, the subseries plot:



This function graphs each month of the series separately. Example: every January together (Ejemplo: todos los meses Enero de la serie, juntos) . Average: blue line.

So, we confirm the fluctuations we saw in the first plot is seasonality every december and january.

Forecast with various methods (considering seasonality).

FIRST METHOD: Use a Benchmark method to forecast. (there are many)

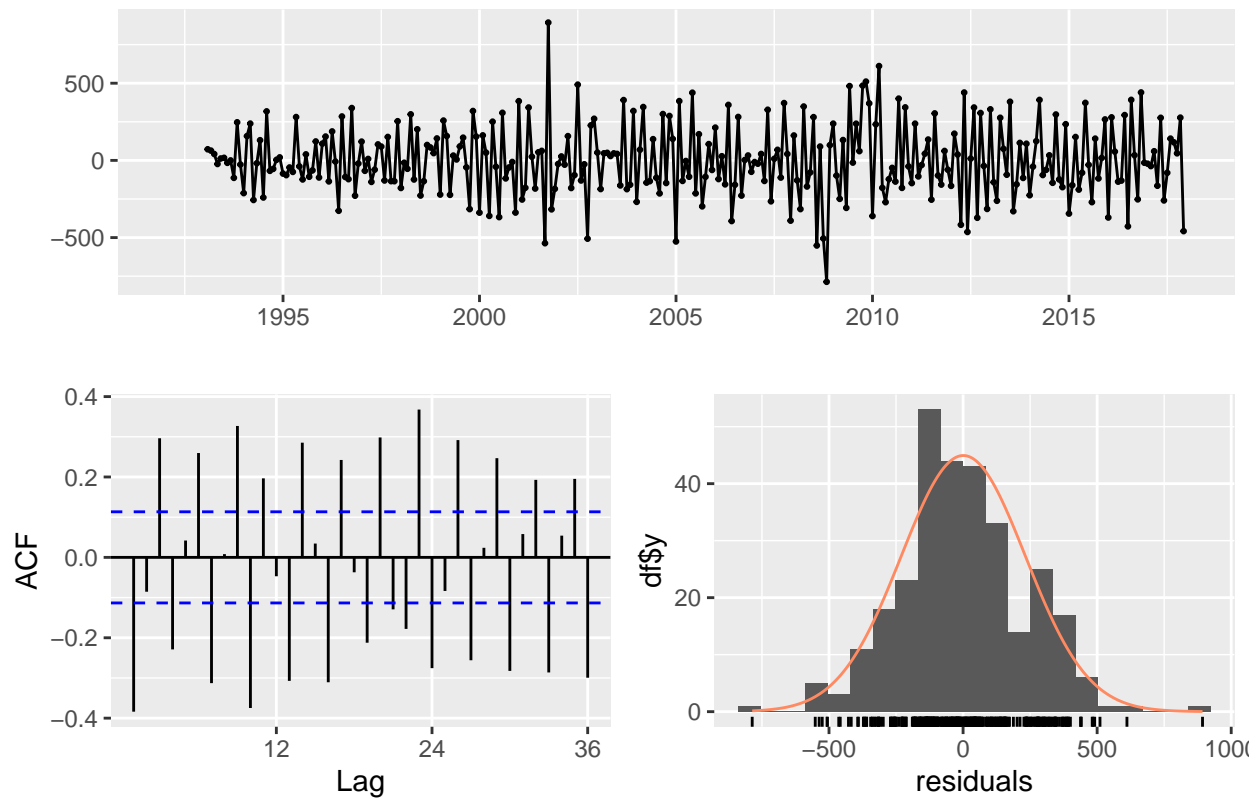
So, lets use the Seasonal naive method as our benchmark: $y_t = y_{t-s} + e_t$

Using the differenced series (1st difference):

```
##
## Forecast method: Seasonal naive method
##
## Model Information:
## Call: snaive(y = DY)
##
## Residual sd: 231.4508
##
## Error measures:
##           ME      RMSE      MAE      MPE      MAPE  MASE      ACF1
## Training set 0.9090517 231.4508 181.9523 27.85462 105.8591    1 -0.3838154
##
## Forecasts:
```

##	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Jan 2018	-3642.61290	-3939.22898	-3345.99682	-4096.24805	-3188.97775
## Feb 2018	1151.01843	854.40235	1447.63451	697.38328	1604.65358
## Mar 2018	588.36867	291.75259	884.98475	134.73352	1042.00382
## Apr 2018	-70.55914	-367.17522	226.05694	-524.19429	383.07601
## May 2018	441.65591	145.03983	738.27199	-11.97924	895.29106
## Jun 2018	85.87742	-210.73866	382.49350	-367.75773	539.51257
## Jul 2018	-653.61935	-950.23543	-357.00327	-1107.25450	-199.98420
## Aug 2018	480.58064	183.96456	777.19672	26.94549	934.21579
## Sep 2018	-177.62796	-474.24404	118.98812	-631.26311	276.00719
## Oct 2018	-274.98494	-571.60102	21.63114	-728.62009	178.65021
## Nov 2018	1254.55161	957.93553	1551.16769	800.91646	1708.18676
## Dec 2018	1293.57742	996.96134	1590.19350	839.94227	1747.21257
## Jan 2019	-3642.61290	-4062.09138	-3223.13442	-4284.14988	-3001.07592
## Feb 2019	1151.01843	731.53995	1570.49691	509.48145	1792.55541
## Mar 2019	588.36867	168.89019	1007.84715	-53.16831	1229.90565
## Apr 2019	-70.55914	-490.03762	348.91934	-712.09612	570.97784
## May 2019	441.65591	22.17743	861.13439	-199.88107	1083.19289
## Jun 2019	85.87742	-333.60106	505.35590	-555.65956	727.41440
## Jul 2019	-653.61935	-1073.09783	-234.14087	-1295.15633	-12.08237
## Aug 2019	480.58064	61.10216	900.05912	-160.95634	1122.11762
## Sep 2019	-177.62796	-597.10644	241.85052	-819.16494	463.90902
## Oct 2019	-274.98494	-694.46342	144.49354	-916.52192	366.55204
## Nov 2019	1254.55161	835.07313	1674.03009	613.01463	1896.08859
## Dec 2019	1293.57742	874.09894	1713.05590	652.04044	1935.11440

Residuals from Seasonal naive method



```
##
## Ljung-Box test
##
## data: Residuals from Seasonal naive method
## Q* = 464.01, df = 24, p-value < 2.2e-16
##
## Model df: 0. Total lags used: 24
```

We see residuals ACF are not ideal, lets look if we can find a better model


SECOND METHOD: Fit ETS method

Another type we can try is Exponential smoothing model.

They (exponential smoothing model) use weighted averages of past observations to forecast new values
Exponential smoothing models are based on a description of the trend and seasonality in the data

We can use regular data (not the differenced one)

```
## ETS(M,A,M)
##
## Call:
## ets(y = Y)
##
## Smoothing parameters:
##   alpha = 0.4528
##   beta  = 2e-04
##   gamma = 0.2343
##
## Initial states:
##   l = 4744.6674
##   b = 35.6917
##   s = 1.2033 1.0226 0.9822 0.9768 0.9912 0.985
##         1.029 1.009 1.003 0.9722 0.9462 0.8794
##
## sigma: 0.0197
##
##      AIC      AICc      BIC
## 5040.721 5042.803 5104.352
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -12.10585 186.3035 140.0386 -0.1532637 1.503134 0.3074385
##              ACF1
## Training set 0.009306847
```



```
##
##  Ljung-Box test
##
## data:  Residuals from ETS(M,A,M)
## Q* = 303.62, df = 8, p-value < 2.2e-16
##
## Model df: 16.    Total lags used: 24
```

Apparently it did not improve, still correlation left in the data and residuals look almost the same But it fits better, since the Residual SD is much smaller.

Lets do another type of model:

THIRD METHOD: Fit an ARIMA Model

Exponential smoothing and ARIMA models are the two most widely used approaches to time series forecasting.

- ARIMA models aim to describe the autocorrelation in the data.
- Needs to be stationary data




So lets use differenced data, there is seasonality.

We can use regular data and tell it to take 1st difference ($d=1$) and 1st seasonality difference ($D=1$).

- “stepwise = FALSE” finds the model that fits better and try every single possibility, so it could take a lot of time
- “aproximation = FALSE” will show the exactly result, not an approximation
- “Trace” will print all the models

```
## Series: Y
## ARIMA(0,1,1)(2,1,2)[12]
##
## Coefficients:
##          ma1      sar1      sar2      sma1      sma2
##      -0.3545  0.8436  -0.7189  -1.1980  0.6266
## s.e.   0.0580  0.0612  0.0826   0.0762  0.0837
##
## sigma^2 = 30589: log likelihood = -1973.05
## AIC=3958.1   AICc=3958.38   BIC=3980.3
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 2.501669 169.7773 126.3304 -0.002525628 1.352898 0.2773436
##              ACF1
## Training set -0.0009506685
```

Residuals from ARIMA(0,1,1)(2,1,2)[12]



```
##
## Ljung-Box test
##
```

```
## data: Residuals from ARIMA(0,1,1)(2,1,2)[12]
## Q* = 64.685, df = 19, p-value = 6.868e-07
##
## Model df: 5. Total lags used: 24
```

```
#SD = sqrt(sigma^2)
sqrt(30589)
```

```
## [1] 174.8971
```


This model fits better, there is not much autorrelation in residuals (ACF), but is not perfect!

It means there is better model for this time series, but lets use this model.

FORECAST

Forecast with ARIMA Model

```
fcst <- forecast(fit_arima, h=24)
autoplot(fcst, include = 60)
```



```
print(summary(fcst))
```

```
##
## Forecast method: ARIMA(0,1,1)(2,1,2)[12]
##
## Model Information:
## Series: Y
## ARIMA(0,1,1)(2,1,2)[12]
##
## Coefficients:
##          ma1      sar1      sar2      sma1      sma2
##      -0.3545  0.8436  -0.7189  -1.1980  0.6266
## s.e.   0.0580  0.0612   0.0826   0.0762  0.0837
##
## sigma^2 = 30589: log likelihood = -1973.05
## AIC=3958.1   AICc=3958.38   BIC=3980.3
##
## Error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 2.501669 169.7773 126.3304 -0.002525628 1.352898 0.2773436
##              ACF1
## Training set -0.0009506685
##
## Forecasts:
##      Point Forecast      Lo 80      Hi 80      Lo 95      Hi 95
## Jan 2018      12735.82 12511.68 12959.96 12393.03 13078.61
## Feb 2018      13909.29 13642.50 14176.07 13501.27 14317.30
## Mar 2018      14411.74 14108.24 14715.23 13947.58 14875.90
## Apr 2018      14413.70 14077.47 14749.92 13899.49 14927.91
## May 2018      14903.98 14537.95 15270.02 14344.18 15463.79
## Jun 2018      14748.42 14354.83 15142.02 14146.47 15350.37
## Jul 2018      14335.54 13916.20 14754.89 13694.21 14976.88
## Aug 2018      14772.59 14328.98 15216.20 14094.15 15451.03
## Sep 2018      14402.36 13935.75 14868.97 13688.74 15115.97
## Oct 2018      14324.98 13836.45 14813.50 13577.84 15072.12
## Nov 2018      15339.64 14830.14 15849.15 14560.42 16118.87
## Dec 2018      16636.27 16106.62 17165.93 15826.24 17446.31
## Jan 2019      13399.32 12795.77 14002.88 12476.26 14322.39
## Feb 2019      14538.64 13889.81 15187.47 13546.35 15530.94
## Mar 2019      14994.55 14303.41 15685.69 13937.55 16051.55
## Apr 2019      15120.03 14389.03 15851.03 14002.06 16238.00
## May 2019      15599.12 14830.32 16367.92 14423.34 16774.90
## Jun 2019      15408.33 14603.50 16213.16 14177.45 16639.21
## Jul 2019      15160.02 14320.71 15999.34 13876.41 16443.64
## Aug 2019      15449.56 14577.13 16321.99 14115.29 16783.83
## Sep 2019      14985.54 14081.20 15889.88 13602.47 16368.61
## Oct 2019      14959.88 14024.72 15895.04 13529.67 16390.08
## Nov 2019      15736.02 14771.02 16701.02 14260.18 17211.86
## Dec 2019      17214.65 16220.71 18208.59 15694.55 18734.75
```

It looks realistic - “h = 24” forecasted period (months) - function Include, means the number of months (historic data)